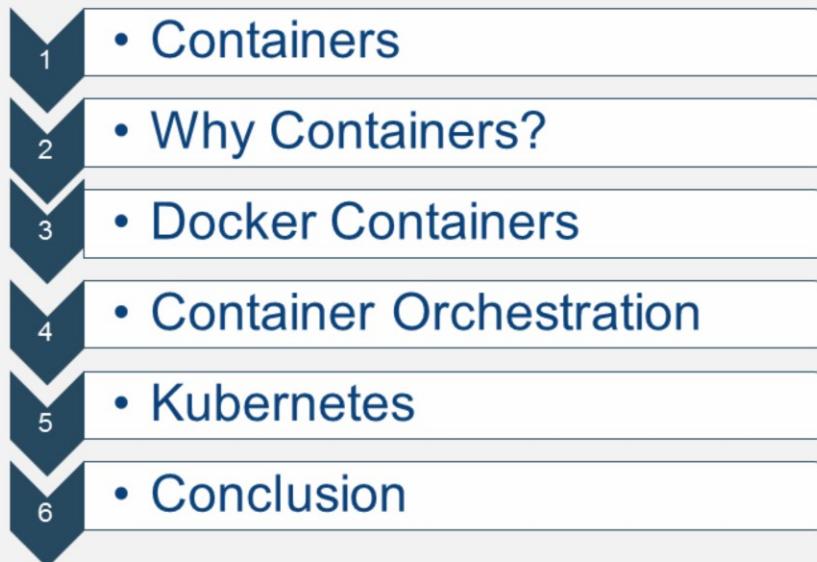


Introduction to Containers

NA Partner Channels – IBM Cloud
Technical Evangelism Team
January 23, 2018

© Copyright IBM Corporation 2017

Agenda



IBM 2017. Internal and Business Partner Use Only

Why Containers?

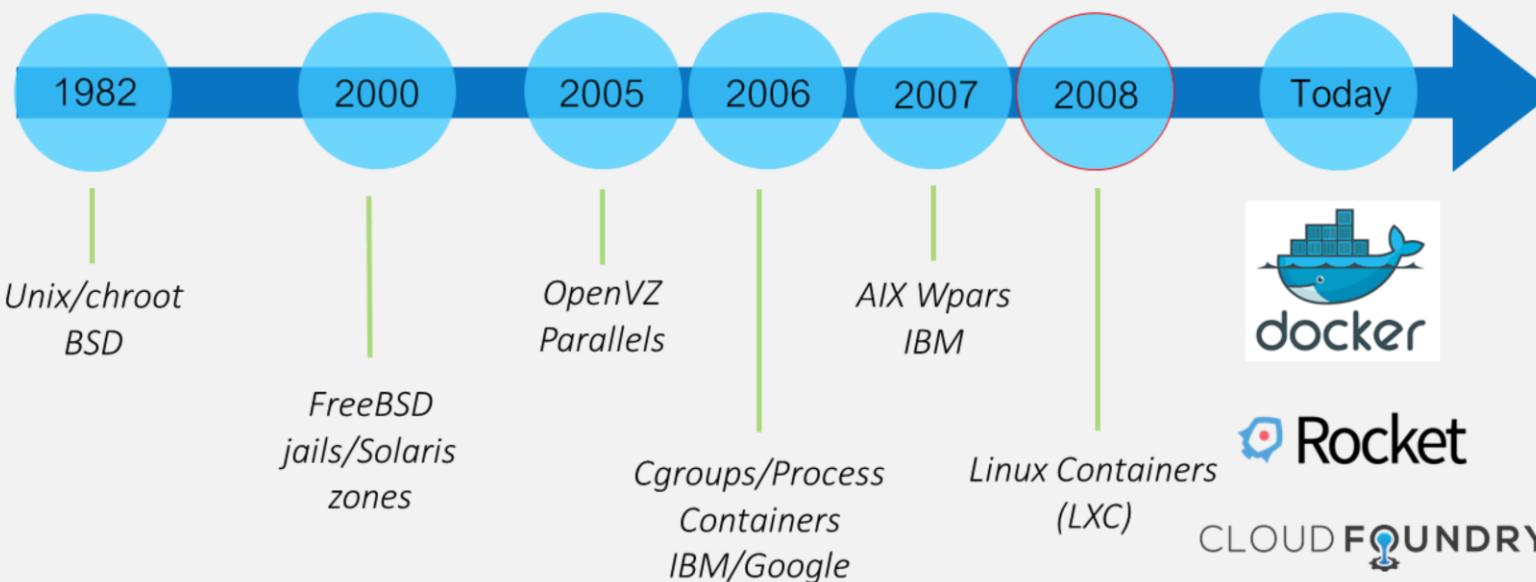
Everyone Loves Containers



- A standard way to package an application and all its dependencies so that it can be moved between environments and run without changes
- Containers work by isolating the differences between applications inside the container so that everything outside the container can be standardized

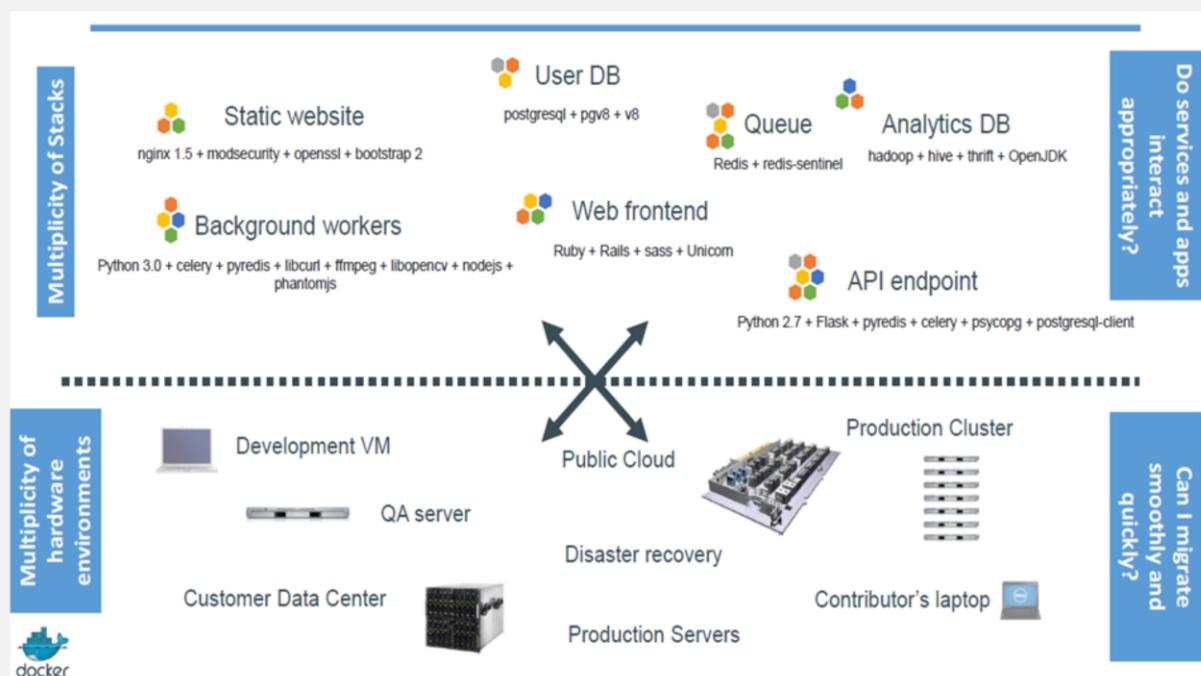
IBM 2017. Internal and Business Partner Use Only

Container History Lesson



IBM 2017 Internal and Business Partner Use Only

The Challenge



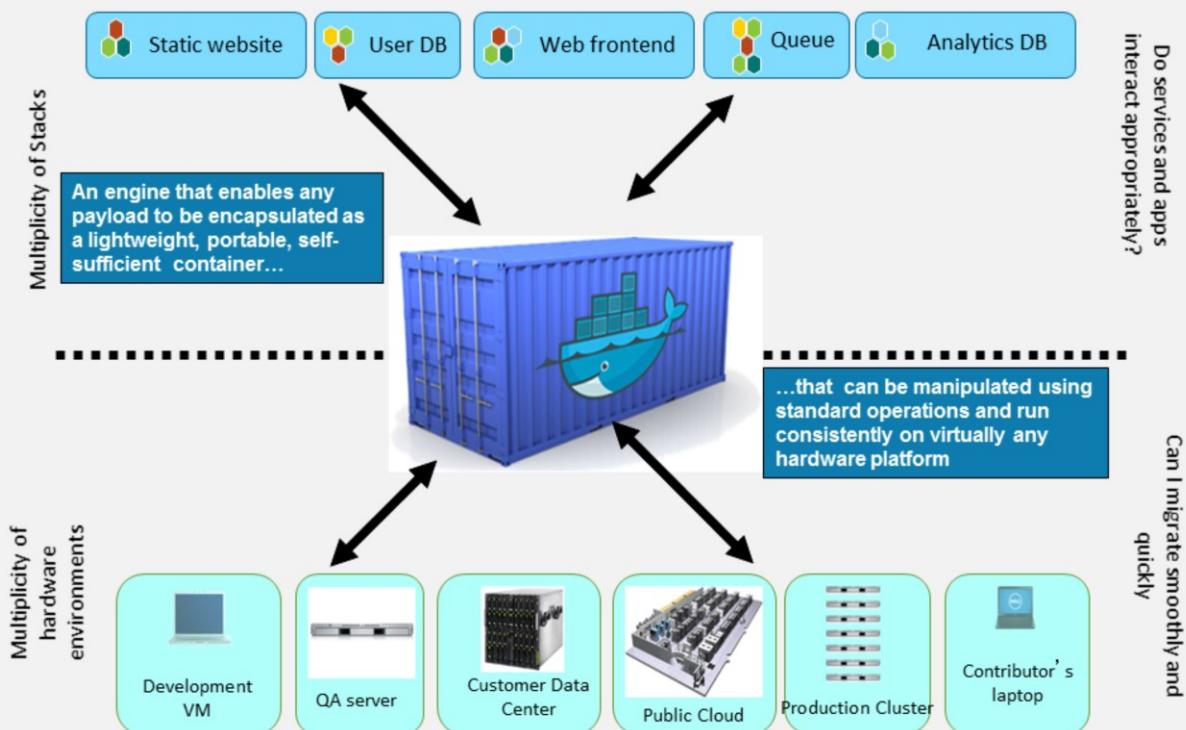
IBM 2017. Internal and Business Partner Use Only

The Matrix from Hell

	Static website	?	?	?	?	?	?	?
	Web frontend	?	?	?	?	?	?	?
	Background workers	?	?	?	?	?	?	?
	User DB	?	?	?	?	?	?	?
	Analytics DB	?	?	?	?	?	?	?
	Queue	?	?	?	?	?	?	?
	Development VM	QA Server	Single Prod Server	Onsite Cluster	Public Cloud	Contributor's laptop	Customer Servers	
								

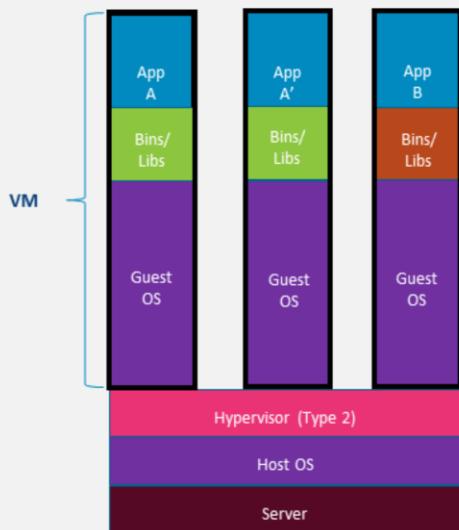
IBM 2017 Internal and Business Partner Use Only

A Shipping Container for Code



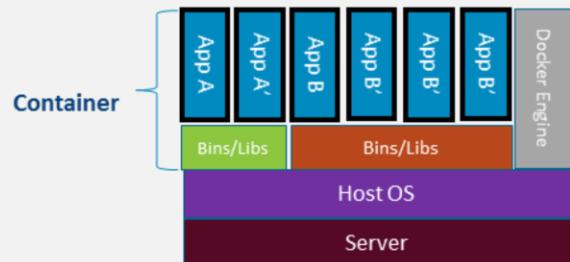
IBM 2017. Internal and Business Partner Use Only

VM's vs. Containers

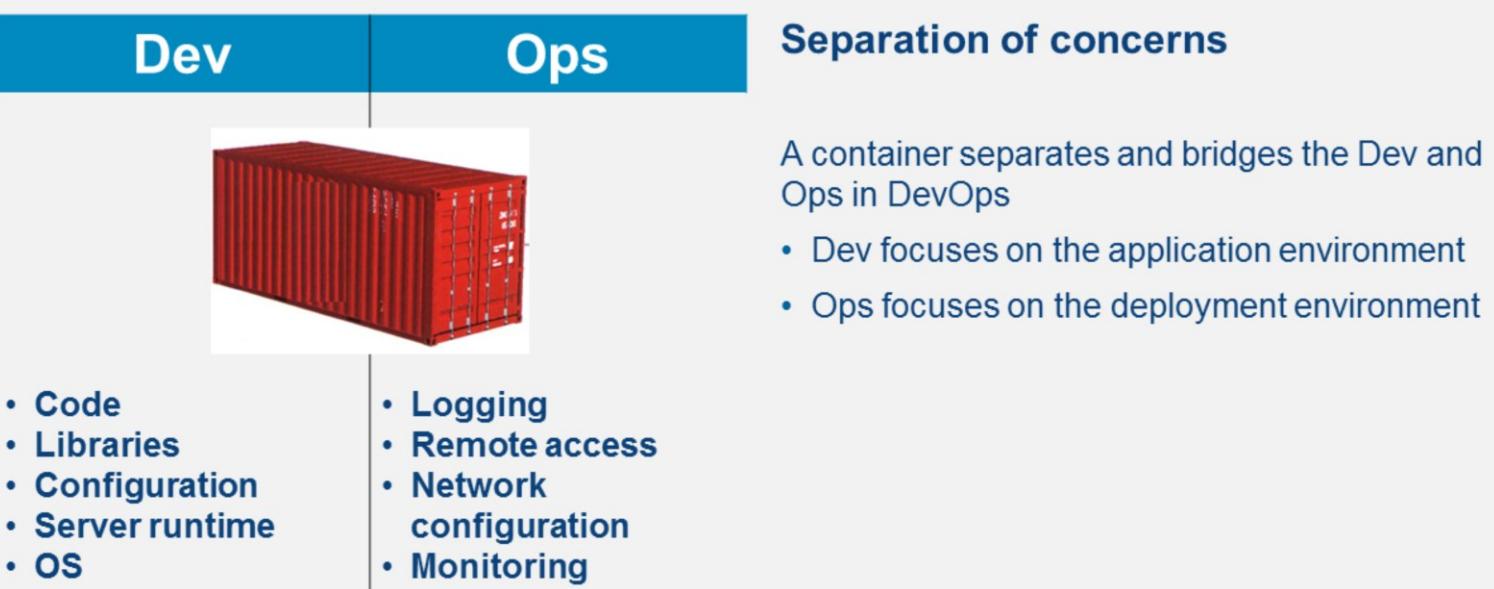


Containers are isolated, but share OS and, where appropriate, bins/libraries

...faster, less overhead



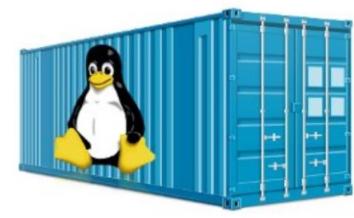
Dev vs. Ops



IBM 2017. Internal and Business Partner Use Only

Container Advantages

- Containers are portable
- Containers are easy to manage
- Containers provide “just enough” isolation
- Containers use hardware more efficiently
- Containers are immutable



© IBM 2017. Internal and Business Partner Use Only

- **Containers are portable**
 - Any platform with a container engine can run containers
- **Containers are easy to manage**
 - Container images are easy to share, download, and delete
 - Especially with Docker registries
 - Container instances are easy to create and delete
 - Each container instance is easy and fast to start and stop
- **Containers provide “just enough” isolation**
 - More lightweight than virtual machines
 - Processes share the operating system kernel but are segregated

- Containers use hardware more efficiently
 - Greater density than virtual machines
 - Especially Docker containers, which can share layers
- Containers are immutable
 - Container images are versions
 - Containers cannot (should not) be patched

Docker Containers

Docker Adoption

Docker enables application development **efficiency**, making deployment more **efficient**, eliminating vendor 'lock-in' with true **portability**



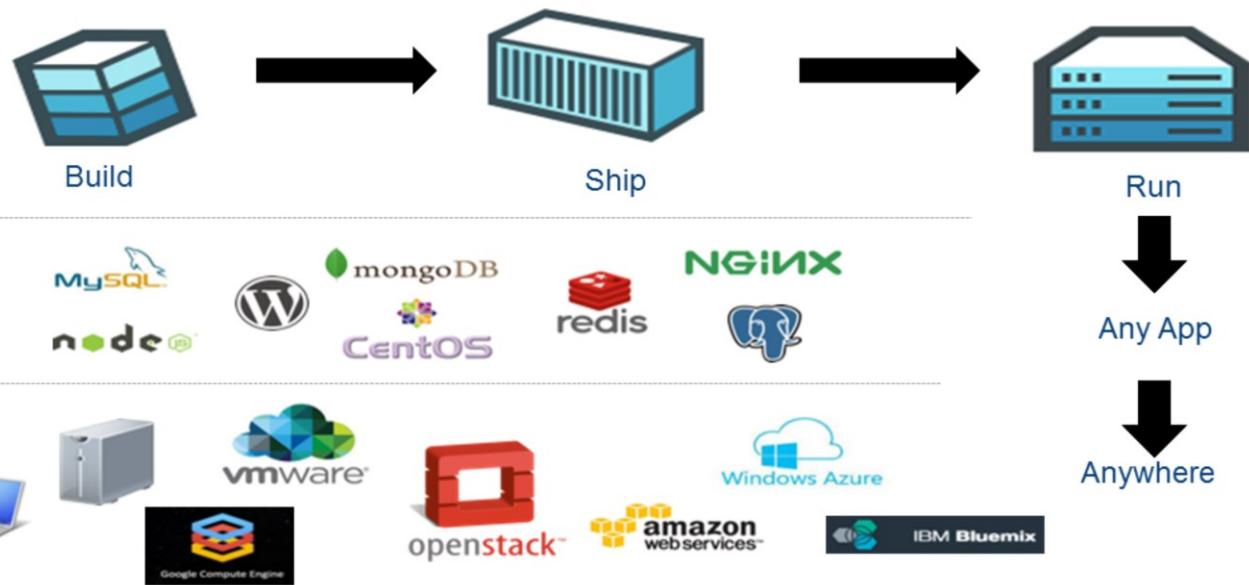
- Open software
- Open contribution
- Open design
- Open governance

IBM 2017 Internal and Business Partner Use Only

- Open software
 - Launched **March 2013**
 - 2.0+ billion downloads of Docker images
- Open contribution
 - 2000+ contributors
 - #2 most popular project
 - 185 community meet-up groups in 58 countries
- Open design
 - Contributors include IBM, Red Hat, Google, Microsoft, VMware, AWS, Rackspace, and others
- Open governance
 - 12 member governance advisory board selected by the community

Docker Mission

Docker is an **open platform** for building distributed applications for **developers** and **system administrators**



IBM 2017. All rights reserved. Confidential - Not for Distribution

Docker Components



Image

- A read-only snapshot of a container stored in Docker Hub to be used as a template for building containers



Container

- The standard unit in which the application service resides or transported

SaaS Enterprise

Docker Hub/Registry

- Available in SaaS or Enterprise to deploy anywhere you choose
- Stores, distributes, and shares container images

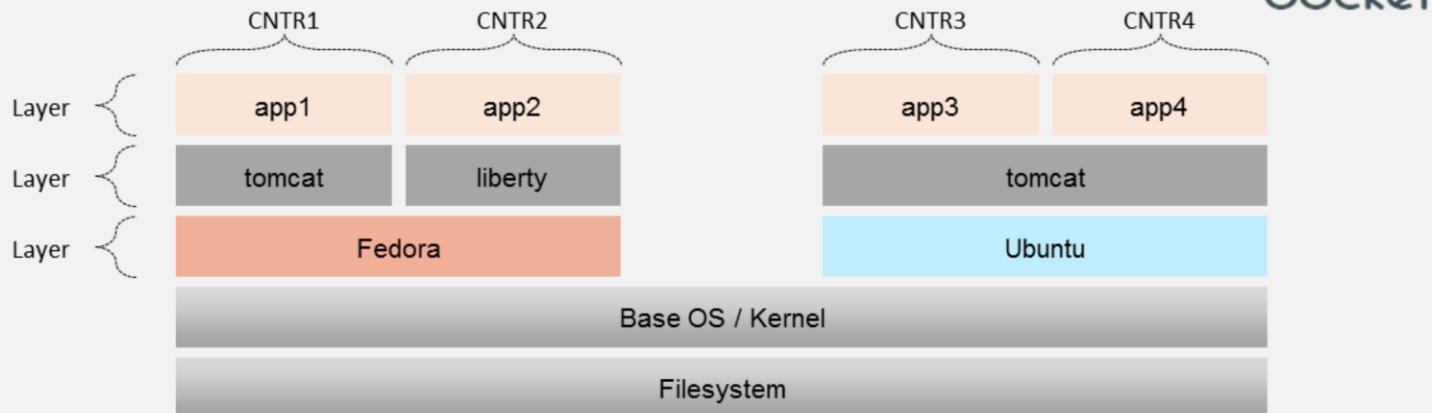


Docker Engine

- A program that creates, ships, and runs application containers
- Runs on any physical and virtual machine or server locally, in private or public cloud
- Client communicates with Engine to execute commands

Docker Containers

- Docker uses a copy-on-write (union) filesystem
- New files (& edits) are only visible to current/above layers



- Layers allow for reuse
 - More containers per host
 - Faster start-up/download time – base layers are "cached"
- Images
 - Tarball of layers (each layer is a tarball)

Containers



Everyone's container journey starts with one container....

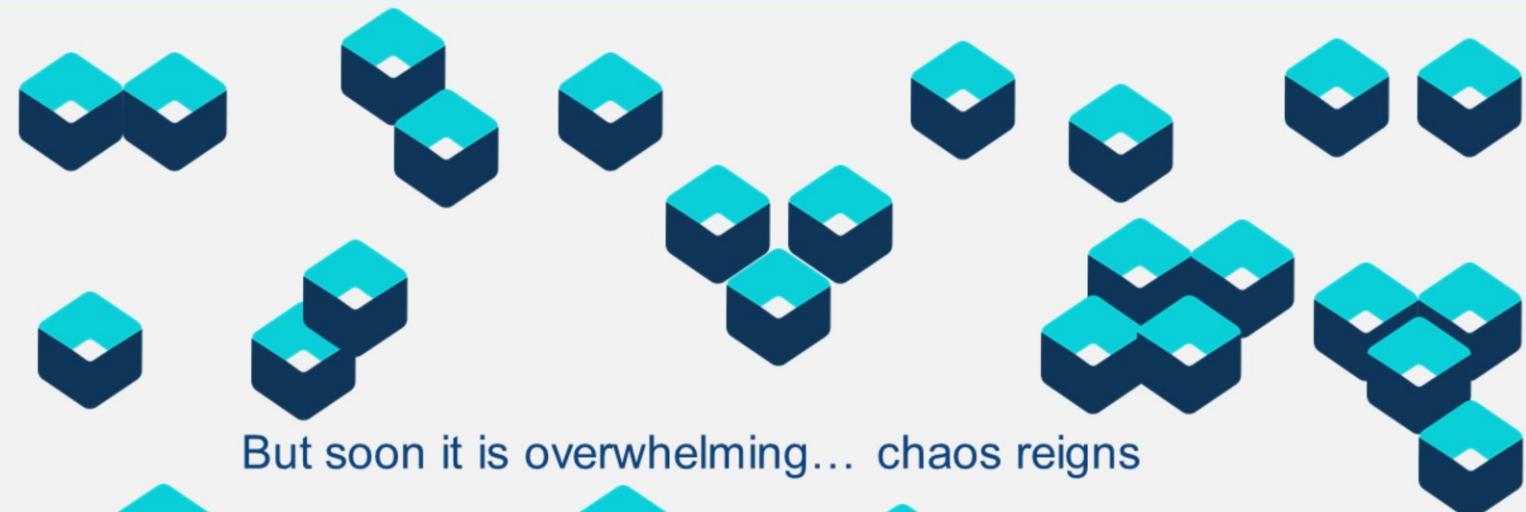
Containers



At first the growth is easy to handle....



Containers



But soon it is overwhelming... chaos reigns



© IBM 2017. Internal and Business Partner Use Only

Orchestration

More to Containers than just Docker

Serverless



PaaS



Container Orchestration



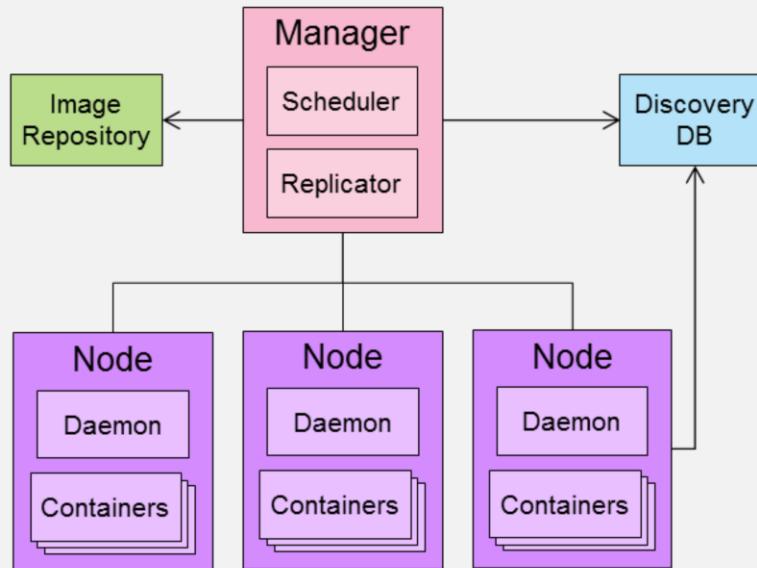
Container Engine



What is Container Orchestration?

- Container orchestration
 - Cluster management
 - Scheduling
 - Service discovery
 - Replication
 - Health management

Container Orchestrator



IBM 2017 Internal and Business Partner Use Only

- Container orchestration
 - Manages the deployment, placement, and lifecycle of workload containers
- Cluster management
 - Federates multiple hosts into one target
- Scheduling
 - Distributes containers across nodes
- Service discovery
 - Knows where the containers are located
 - Distributes client requests across the containers
- Replication
 - Ensures the right number of nodes and containers

- Health management
 - Replaces unhealthy containers and nodes

Container Ecosystem Layers

Layer 6

**Development Workflow
Opinionated Containers**

Layer 5

**Orchestration/Scheduling
Service Model**

Layer 4

Container Engine

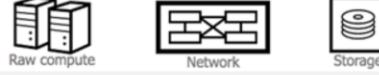
Layer 3

Operating System

Layer 2

Virtual Infrastructure

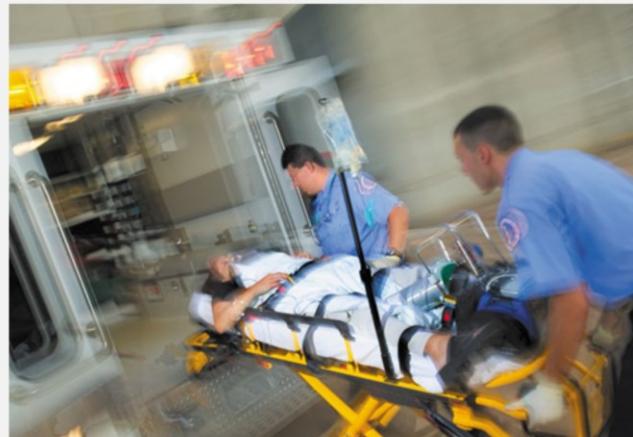
Layer 1

Physical Infrastructure

Kubernetes

What is Kubernetes?

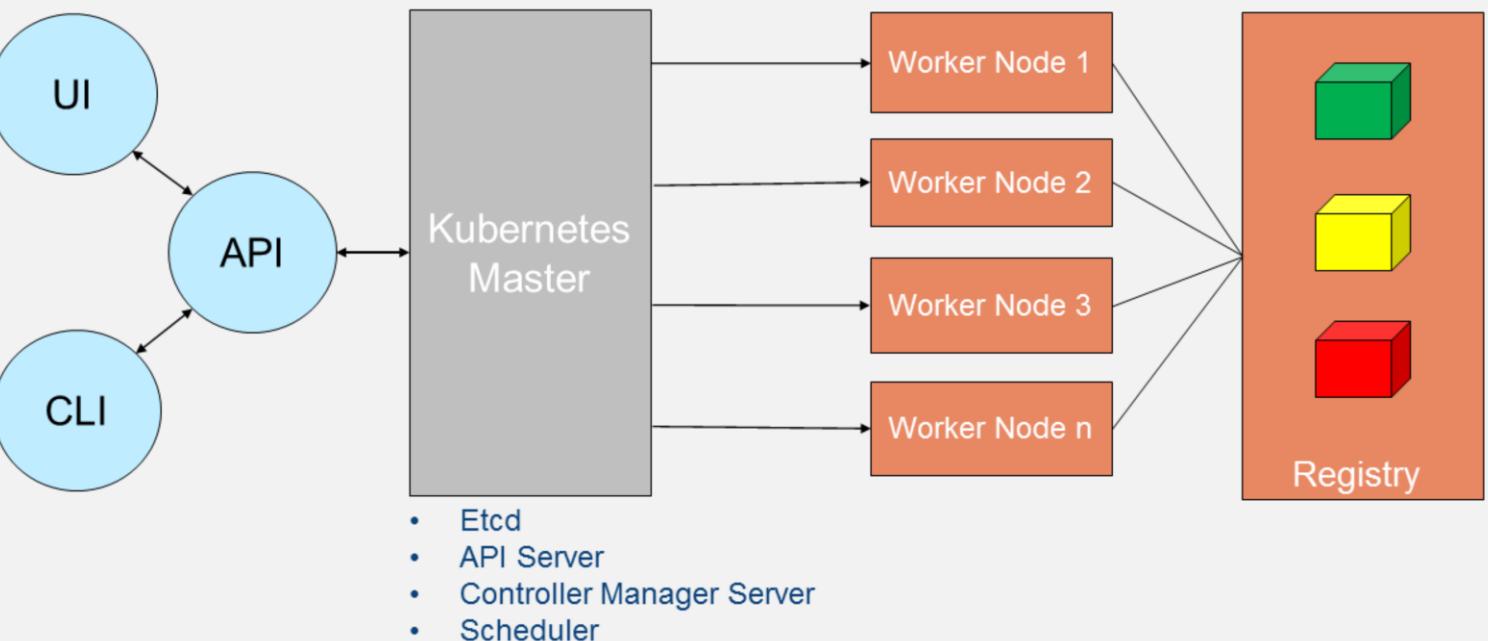
- Container orchestrator
- Manage applications, not machines
- Designed for extensibility
- Open source project managed by the Linux Foundation



- Container orchestrator
 - Runs and manages containers
 - Unified API for deploying web applications, batch jobs, and databases
 - Maintains and tracks the global view of the cluster
 - Supports multiple cloud and bare-metal environments
- Manage applications, not machines
 - Rolling updates, canary deploys, and blue-green deployments
- Designed for extensibility
 - Rich ecosystem of plug-ins for scheduling, storage, networking

- Open source project managed by the Linux Foundation
 - Inspired and informed by Google's experiences and internal systems
 - 100% open source, written in Go

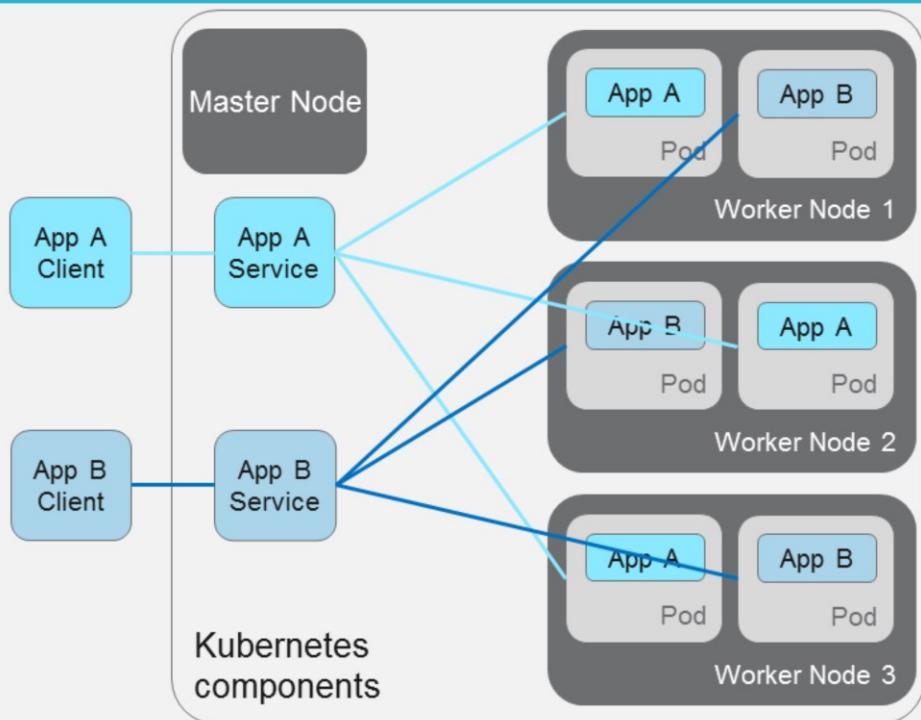
Kubernetes Architecture



© IBM 2017. Internal and Business Partner Use Only

Kubernetes Architecture: Workloads

- Container
 - Packaging of an app
- Pod
 - Unit of deployment
- Service
 - Fixed endpoint for 1+ pods



IBM 2017. Internal and Business Partner Use Only

Kubernetes Terminology: Deployment

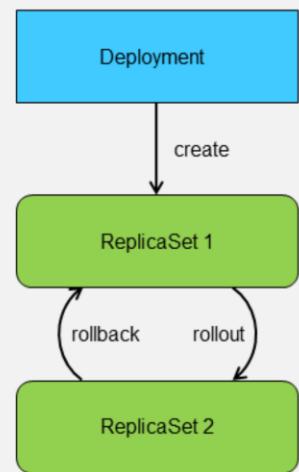
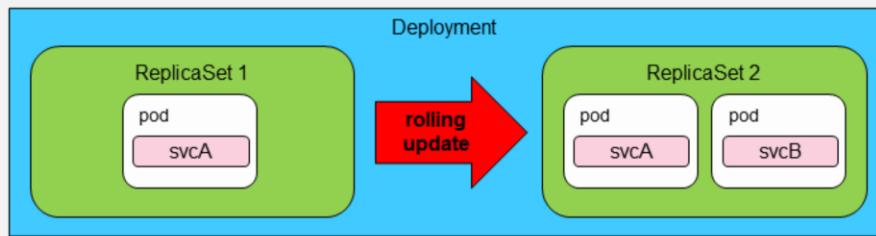


- Deployment

- A set of pods to be deployed together, such as an application
- Declarative: Revising a Deployment creates a ReplicaSet describing the desired state
- Rollout: Deployment controller changes the actual state to the desired state at a controlled rate
- Rollback: Each Deployment revision can be rolled back
- Scale and autoscale: A Deployment can be scaled

- ReplicaSet

- The next-generation ReplicationController
- A set of pod templates that describe a set of pod replicas
- Uses a template that describes specifically what each pod should contain
- Ensures that a specified number of pod replicas are running at any given time

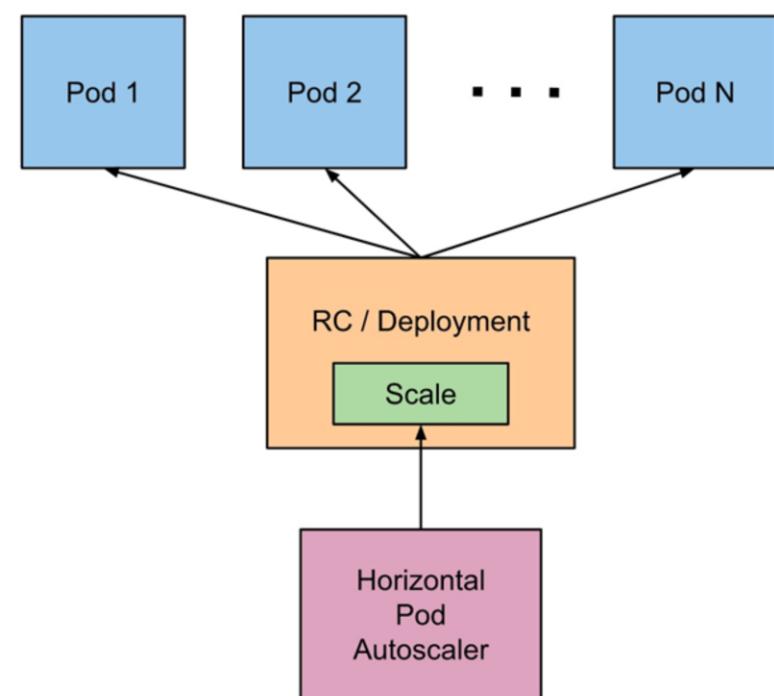


IBM 2017 Internal and Business Partner Use Only

Kubernetes Terminology: Autoscaling

- Horizontal Pod Autoscaling (HPA)

```
$ kubectl autoscale deployment <deployment-name> --cpu-percent=50
  --min=1 --max=10 deployment
" <hpa-name>" autoscaled
```



© IBM 2017. Internal and Business Partner Use Only

- Horizontal Pod Autoscaling (HPA)
 - Automatically scales the number of pods in a replication controller, deployment, or replica set
 - Matches the observed average CPU utilization to the specified target
 - Fetches metrics in two different ways: direct Heapster access and REST client access
 - Kubernetes Heapster enables container cluster monitoring and performance analysis
 - Default config: query every 30 sec, maintain 10% tolerance, wait 3 min after scale-up, wait 5 min after scale-down

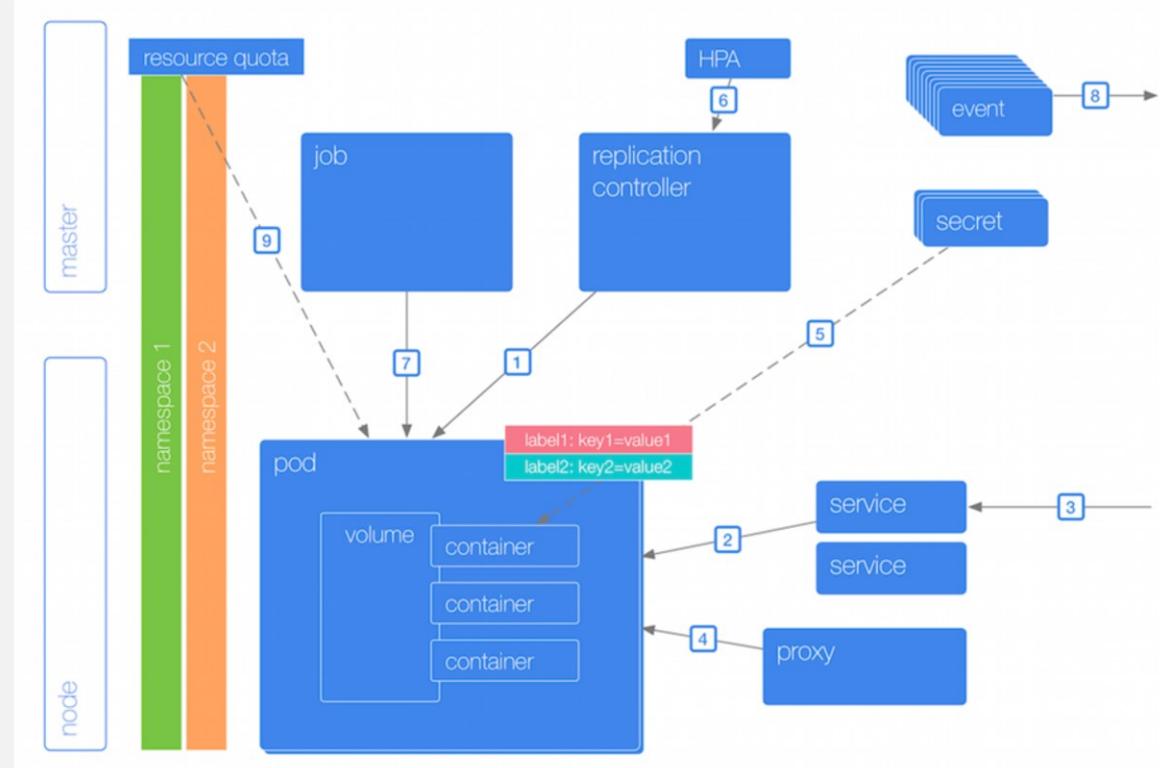
```
$ kubectl autoscale deployment <deployment-name> --cpu-percent=50
```

```
--min=1 --max=10 deployment "<hpa-name>" autoscaled
```

- Creates a horizontal pod autoscaler
 - An HPA instance
 - Maintains between 1 and 10 replicas of the pods controlled by the deployment
 - Maintains an average CPU utilization across all pods of 50%

Kubernetes Terminology: Naming

- Name
- Namespace
- Resource Quota
- Label
- Selector
- ConfigMap
- Secrets



- **Name**
 - Each resource object by type has a unique name
- **Namespace**
 - **Resource isolation:** Each namespace is a virtual cluster within the physical cluster
 - Resource objects are scoped within namespaces
 - Low-level resources are not in namespaces: nodes, persistent volumes, and namespaces themselves
 - Names of resources need to be unique within a namespace, but not across namespaces
 - **Resource quotas:** Namespaces can divide cluster resources
 - **Initial namespaces**
 - `default` – The default namespace for objects with no other namespace

- `kube-system` – The namespace for objects created by the Kubernetes system

- **Resource Quota**

- Limits resource consumption per namespace
- Limit can be number of resource objects by type (pods, services, etc.)
- Limit can be total amount of compute resources (CPU, memory, etc.)
- Overcommit is allowed; contention is handled on a first-come, first-served basis

- **Label**

- Metadata assigned to Kubernetes resources (pods, services, etc.)
- Key-value pairs for identification
- Critical to Kubernetes as it relies on querying the cluster for resources that have certain labels

- **Selector**

- An expression that matches labels to identify related resources

- **ConfigMap**

- Configuration values to be used by containers in a pod
- Stores configuration outside of the container image, making containers more reusable

- **Secrets**
 - Sensitive info that containers need to read or consume
 - Encrypted in special volumes mounted automatically

Common Kubernetes Commands

- Get the state of your cluster
\$ kubectl cluster-info
- Get all the nodes of your cluster
\$ kubectl get nodes -o wide
- Get info about the pods of your cluster
\$ kubectl get pods -o wide
- Get info about the replication controllers of your cluster
\$ kubectl get rc -o wide
- Get info about the services of your cluster
\$ kubectl get services
- Get full config info about a Service
\$ kubectl get service NAME_OF_SERVICE -o json
- Get the IP of a Pod
\$ kubectl get pod NAME_OF_POD --template={{.status.podIP}}
- Delete a Pod
\$ kubectl delete pod NAME
- Delete a Service
\$ kubectl delete service NAME_OF_SERVICE

Helm & Helm Chart

- Helm
 - Package manager for Kubernetes
 - Used to manage Kubernetes applications
- Helm Chart
 - Used to define, install, and upgrade complex Kubernetes applications
 - Easy to create, version, share and publish
 - Expressed in “Yet Another Markup Language” (YAML) files



Conclusion

Conclusion

- Why containers?
- Docker containers
- Container orchestration
- Kubernetes

IBM 2017 Internal and Business Partner Use Only

Resources

- Docker tutorial
 - <https://docs.docker.com/get-started/>
- Kubernetes tutorial
 - <https://kubernetes.io/docs/tutorials/kubernetes-basics/>
- The Evolution of Linux Containers and Their Future
 - <https://dzone.com/articles/evolution-of-linux-containers-future>
- Introduction to container orchestration
 - <https://www.exoscale.ch/syslog/2016/07/26/container-orch/>
- TNS Research: The Present State of Container Orchestration
 - <https://thenewstack.io/tns-research-present-state-container-orchestration/>
- Large-scale cluster management at Google with Borg
 - <https://research.google.com/pubs/pub43438.html>

IBM 2017 Internal and Business Partner Use Only



IBM 2017. Internal and Business Partner Use Only